

REU Tutorial 2019

Abid Khan, Sean Connelly, Cunwei Fan, and Patchara Wongsutthikoson
Edited: Eric Connelly, Imran Sultan, Minh Nguyen

March 10, 2019

1 Introduction

This tutorial will familiarize you with the basic software and methods used in this REU. It is not comprehensive, but upon completion, you should be able to contribute effectively to the group. Previous incoming students have typically completed this tutorial in a few weeks.

2 Accounts and Keys

One of the senior members will join you to get this section done. You can take your time, but getting these done quickly will make your life easier.

- Get keys for Loomis and the REU office.
- Talk to Rebecca Wiltfong (rwiltfon@illinois.edu, 290L LLP) about permissions for the REU webpage.
- Get your Blue Waters Token and have the account activated. You should have gotten an email about it. Talk to Milton Ruiz (setisam@gmail.com) if you come across troubles.

3 Linux

This group uses Linux exclusively. Over the years, you will learn to use it and eventually prefer it over any other operating system. One of the most important and powerful tools you will need to master on Linux is the **terminal**. How to use terminal will become clear as you work with it. A good way to start is by familiarizing yourselves with the following basic commands (i.e. by googling, etc.): `cd`, `ls`, `mkdir`, `touch`, `cp`, `rm` (be *very careful* with this one). A quick way to learn what a command does is to type `man commandname` into the terminal.

After that, you can look at some more useful commands below.

grep This command outputs items that match your input pattern strings. It is very useful in finding files and filter the output of another command when used with pipe (`|`). eg: `grep 1200 foo.txt` will output the lines that has "1200" in file `foo.txt`.

pipe (`|`) This command makes compound commands available, allowing you to perform more complicated operations in one command line. It will feed the output of previous command as the input for the next one. eg: `ls | grep foo`, this use the output of `ls` command as the input for `grep` command. This command line will display the items in current directory that has pattern string "foo" in it.

redirect (`>`) This command will redirect the output display to another proper place (we usually use this to output into a text file.) eg `echo hello`

`world > new.txt` will modify the file "new.txt" (or make a new file "new.txt", if it doesn't exist) to have "hello world" as its content. Note that `>` erases the file before writing, while `>>` appends to it.

sort What do you think it does?

uniq Delete *consecutive* duplicate lines of a file. See **awk** for how to delete *non-consecutive* duplicate lines.

scp At times we will be needing to send files or data to each other for collaboration. How do we do that? EMAIL?! Where's the fun in that? To send a file from one machine to another type

`scp filename colten1@machine.physics.illinois.edu:/home/colten1/Desktop,`
where instead of "machine" put in the name of the computer you want to send it to. Also you can retrieve a file from another machine using the same command:

`scp colten1@machine.physics.illinois.edu:/full/path/to/file /path/to/folder`
If you want to transfer a directory instead, put an `-r` after `scp`

scripting Bash supports a few simple common functions like for loop and if condition. These function allow you to write more complicated scripts that can perform a series of operations. I'm not going to write down the usage of those function, as most of them you should easily find documentations on google whenever you need them. The only one thing I want mention is that in bash scripts, when you want to refer to a variable you defined earlier, put a dollar sign (\$) in the front. eg. `var=1; echo $var`

Vim Get familiar with a terminal-based text editor, such as Vim (some alternatives are Emacs and nano). This will come in very handy for quickly reading/editing files. Vim is quite advanced and there's no need to understand everything (*yet*), but do learn the basics, such as editing and saving files (and exiting out of Vim, of course).

awk (optional) Technically a programming language, but useful for editing text files. Ubiquitous in Unix-like systems. This command allows you to perform data extraction based on patterns. A few simple ways of using it will be `awk 'print $2'` (print 2nd column of a file), and `awk '!a[$0]++'` (remove non-consecutive duplicate lines)
Refer to <http://www.pement.org/awk/awk1line.txt>

sed (optional) Short for 'stream editor'. This is a very powerful string editing tool which allows you to perform text deleting and modifying operations. Refer to <http://sed.sourceforge.net/sed1line.txt>

3.1 Exercises

Don't spend too much time trying to work these out yourselves; if you get stuck, feel free to use Google to whenever you need to. These are mainly designed for reference.

1. Use `wc -l` and `ls` to count the number of files in a folder.
2. Use terminal to remove the extension part of the file name. eg. "foo.txt" \rightarrow "foo".
3. Use the `head` and `tail` command to remove the first and last line of a text file.
4. Use `grep` command to delete all the lines in a text file that contains a specific pattern (say "NaN"). *Hint: -v*
5. Use `sort` command to sort a text file based on the value in the second column in each line. Assume your text file has multiple rows and columns. *Hint: -k, -n*

4 Basic Imaging

The group presently uses VisIt (<https://wci.llnl.gov/codes/visit/>) to render movies and still images. This software is installed on all REU machines, as well as Blue Waters. You are encouraged to read the manuals at:

<https://wci.llnl.gov/codes/visit/manuals.html>

To begin using VisIt, simply type 'visit' into a terminal of your choice.

4.1 Exercises:

1. Useful to us is the `.3d` data format. The `.3d` data format just contains four columns of the form `x y z scalar`. Create a `.3d` datafile of the scalar field $d = x^2y^3/z^5$ and visualize it in VisIt using three unique colormaps. Be creative and try to explore VisIt's features.
2. Next, create and visualize a sphere, using Python, C++ or whatever language you want, with radius 2 at a distance 8 from the origin. *Hint: the Delaunay operator creates a surface from a set of points.*¹ Display the sphere in VisIt.
3. You will be given a VTK file representing a vector field. If you don't have it, ask somebody in the group. Visualize the vector field contained in this VTK file. Create seed particles for streamlines and visualize these as well.

5 Advanced Imaging

The goal of this section is to get you to use the skills and knowledge you have acquired so far. You will do this by visualizing the magnetic field of a current loop in two different ways. Upon successful completion of this exercise you should have the necessary background to start learning the daily work of the group.

¹The Delaunay operator is not enabled in VisIt by default. Turn it on by going to Options > Plugin manager, making sure to save settings. You'll have to restart VisIt.

5.1 Exercise Instructions:

1. First you should generate the x , y , and z components of the magnetic field of a loop of current. The loop has a current of 4A and an area of $4\pi m^2$. Generate this data for an area of $12m \times 12m \times 12m$. The resolution choice is up to you. However, you will be integrating over this data so choose wisely.
 - Useful equations can be found: http://en.wikipedia.org/wiki/Magnetic_moment
 - https://en.wikipedia.org/wiki/Biot-Savart_law
 - Griffiths E+M textbook.
2. Your next task is to visualize this field in VisIt.
3. Following that, generate a data file containing a collection of positions corresponding to points on the ring of current. You shall then use these points as seeding points for streamlines.
4. Your next task is to generate your own field lines and to visualize them in VisIt. You will use the seed points from above as starting points for each field line. For the integration, please implement the fourth order Runge-Kutta method. You can find it on Wikipedia.
5. Compare your field lines with those generated by VisIt. Then, just play around with the seed points until you start to get a feel for them.
6. Save a high quality image of the field and stream lines.

6 Gravitational Waves

Just as how E-M waves come in two polarizations, so do gravitational waves, namely h_+ and h_\times . You will be generating movies of these two polarizations for various systems. The data that is obtained from the post-docs to generate gravitational wave movies is called `Psi4_rad.mon.3`. This file contains columns of data. The first column is the time column, and the columns after that are coefficients ψ_4^{lm} in the equation

$$\psi_4(t, r, \theta, \phi) = \sum_{l=2}^{\infty} \sum_{m=-l}^l \psi_4^{lm}(t, r) \cdot {}_{-2}Y_{lm}(\theta, \phi), \quad (1)$$

where the ψ_4^{lm} columns are evaluated at an extraction radius R . In equation 1, ψ_4 is decomposed into $s = -2$ spin-weighted spherical harmonics. From equation 1, we can obtain h_+ and h_\times by

$$\psi_4 = \ddot{h}_+ - i\ddot{h}_\times \quad (2)$$

Our code, "hplus_hcross.cpp", takes ψ_4^{lm} , the ADM mass M , and the extraction radius R as inputs. It outputs h_+ and h_\times , from which we take and

generate our movies. The code does this by fast Fourier transforming the equation

$$\sum_{l=2}^{\infty} \sum_{m=-l}^l \psi_4^{lm}(t, r) \cdot {}_{-2}Y_{lm}(\theta, \phi) = \ddot{h}_+ - i\ddot{h}_\times, \quad (3)$$

which is just a combining of equations 1 and 2. It then does a backward fast Fourier transform to obtain coefficients C_{lm} in the equation

$$\frac{1}{2}(h_+ - ih_\times) = \sum_{l,m} C_{lm} \cdot {}_{-2}Y_{lm}, \quad (4)$$

from which we can obtain h_+ and h_\times . The ${}_{-2}Y_{2,2}$ and ${}_{-2}Y_{2,-2}$ modes dominate this sum, so we can approximate h_+ and h_\times as

$$h_+ \approx 2\Re(C_{2,2} \cdot {}_{-2}Y_{2,2} + C_{2,-2} \cdot {}_{-2}Y_{2,-2}) \quad (5)$$

$$h_\times \approx -2\Im(C_{2,2} \cdot {}_{-2}Y_{2,2} + C_{2,-2} \cdot {}_{-2}Y_{2,-2}) \quad (6)$$

6.1 Calculating h_+ and h_\times

Zach's code is called to sum over the ${}_{-2}Y_{lm}$ modes and evaluate h_+ and h_\times at a chosen θ and ϕ . It is assumed that in the far zone the amplitudes of h_+ and h_\times drop off like $\frac{1}{R}$. Therefore, Zach's code removes the radial dependence and gives $R \cdot h_+$ and $R \cdot h_\times$. The output from Zach's code when it is passed a chosen (θ_i, ϕ_i) is $R \cdot h'_+(\theta_i, \phi_i)$ and $R \cdot h'_\times(\theta_i, \phi_i)$ evaluated at a series of values of $\frac{t}{M}$, each separated from the next by some timestep Δt .

At a given point (r_j, θ_i, ϕ_i) (defined in a coordinate system with origin at the center of mass of the initial configuration) and time t_k , the amplitudes of h_+ and h_\times will be

$$h_+ = \frac{h'_+(\theta_i, \phi_i)(t')}{r_j}$$

$$h_\times = \frac{h'_\times(\theta_i, \phi_i)(t')}{r_j}$$

where h'_+ and h'_\times are the functions obtained from Zach's code and t' is the retarded time $t' = t - \frac{r_j}{c}$ (c being the speed of light, (taken to be $c = 1$)). Since t only comes in discrete intervals, Δt , we choose radii $r_j = n\Delta t$ and times $t_k = m\Delta t$ for integer n, m at which to evaluate the amplitude of the wave. This ensures that the retarded time $t' = t_k - r_j = (m - n)\Delta t$ is an integer multiple of Δt and can be evaluated with the data given by Zach's code. If t' is less than zero then the amplitude of the wave is taken to be zero, because this implies that $r_j > ct$ (i.e. the wave has not yet reached this radius).

6.2 Data-Processing:

1. In order to make a gravity wave movie, you will be given a number of raw Psi4 datasets in the form of Psi4.rad.mon.#, together with black hole event horizon data and density data. You may also have a diagnostic datafile that gives you the ADM mass, extraction radius etc. Make sure you have all these parts before proceeding.
2. Next, choose which Psi4 file to work with. Generally speaking, we prefer to use a larger extraction radius since waveforms generated by numerical codes at smaller radius around black hole are less trustworthy. However, the problem is sometimes waveforms with larger extraction radius have abnormal behaviors, thus you need to check the Psi4 data yourself. One way to check the data is using gnuplot. For example, you can type: `p "/Psi4.rad.mon.#"` using 1:2 w l.
3. Due to checkpointing, the raw data has a lot of duplicated timesteps. Duplicate timesteps, and those containing a 'NaN' datum must be removed ere proceeding. *Have a well-sorted data or later codes will either crash or give nonsense.* (This exercise is not trivial, because you may run into some tricky problems.)
4. With the sorted data, you can proceed with our main processing code: `hplus_hcross.cpp`. The code will read-in the raw Psi4 data in certain radius, calculate the waveforms, and output to time-series VTK files.
5. There are a couple of things you need to check in the code. In the function `Write3D`, you will need to specify the ADM mass and the extraction radius for example. Then you will need to specify the spatial dimensions (usually we use 150*150) and spatial resolution (use an odd number to avoid singularity at $z=0^2$). Make sure the spatial resolution is neither too low nor too high because then the final data will be unmanagably huge (just gives you an idea, 151*151 may works fine).
6. Compile the code with library files `DataFile.cpp` and `DataFile.h` (always check out codes at `Allo:/home/projects/codes`). Refer to the commentary to compile.
7. It will some time for the code to run.

6.3 Subtle Points

Notice, From equation 2, we see that the second derivative of h_+ or h_\times is Ψ . If we use notation $c = G = 1$, we may see that mass M , radius r and time t has the same units (as $[r] = c[t]$ and $G[M]/c^2 = [r]$). Thus, we see that Ψ_4 has dimension as $[t]^{-2}$ instead of dimensionless. To deal with dimensional quantity, we usually make it dimensionless. Thus, the code "`hplus_hcross.cpp`" assumes

²This may no longer be true. I don't know.

the Ψ_4 files it processes has no dimension. To make the Ψ_4 files dimensionless, we need to rescale the Ψ_4 with a length scale $K^{n/2}$. Where if you know polytropic equation of state, K is the constant in

$$P = K\rho^\Gamma, \quad \text{and} \quad \Gamma = 1 + \frac{1}{n} \quad (7)$$

Thus, for the first column of "Psi4.mon.rad.#" we need to divide $K^{n/2}$ and for the remaining columns we need to multiply by $(K^{n/2})^2 = K^n$. As the first column is time which has unit of length and the remaining columns are Ψ_4 which has units of $[r]^{-2}$. After the "Psi4.rad.mon.#" is rescaled to dimensionless, you can put it into the "hplus_hcross.cpp" to generate waveforms. (Another suggestion, you can rescale it after you sort the Psi4 file "Psi4.rad.mon.#".) However, for pure black hole cases, we may not have polytropic gas and thus we just use $K = 1$. For neutron star cases or other cases that you may need polytropic gas, please ask the post docs about the constants you want.

Another subtle point is that, if you remember, the waveforms h_+ and h_\times are generated from Ψ_4 by backward fast Fourier transform. Also, the Ψ_4 for magnetic field cases may not start from $t = 0$ but start from some positive t_0 . Thus, in the backward fast Fourier transform which is some kind of integration, we may lose some information from $t = 0$ to $t = t_0$. However, if the wave last for a long time, the lost information may not play a big role. If the wave last for very short time, you may ask the post docs or Ph.Ds for the data between $t = 0$ to $t = t_0$. Thus, to determine whether the part of information is crucial, you may want to generate the 1D waveform first and then check whether the plot is wave like and makes sense. If everything is fine, you can keep going; but if the 1D plot does not make sense, you may need the lost part of information. (But usually, the given file has enough information. If you find your 1D plot problematic, please check first whether everything is scaled correctly and sorted finely.)

6.4 Superimposing density profiles

You may be asked to create a film where the density profile is superimposed over the gravitational radiation waveform. Unfortunately, VisIt can only render one volume plot via raycasting. The following method should circumvent this limitation.

1. Film a density movie with the same camera settings as your GW movie
2. Add the script 'batch-set-alpha.scm' to `/.gimp-BLAHBLAHVERSIONNUMBER/scripts/`
This script sets a color to be transparent, and can be used to remove the background from the density images.
3. Make sure the script is set to remove the background color.
4. run `gimp -i -b '(batch-set-alpha "FILENAME") -b 'gimp-quit 0)` on the density files

5. use `imagemagick (composite)` to superimpose the density images onto the GW images

Besides the direct command way, we also have a developed python code to work in the `gimp python fu console`.

It is entirely likely there is a better way to accomplish this, but we haven't bothered to think of it.

7 Python Scripting

VisIt can be scripted via python, which allows easy standardization of movie parameters. The scripting is relatively straightforward. Ask a senior member to provide a preexisting script, then adapt this script to automatically create a movie of the gravity wave you've been examining. Try adjusting the camera settings to zoom and rotate around the waveform.

Python Scripting: You are going to make a movie that will be filmed entirely using a python script. That means the data will be loaded, the plots will be set up, and your dataset will be iterated through time. Also the camera should move around all pretty and such. I'll have you create this script piece by piece so that its all not a giant mess.

You are going to need a time varying vector dataset. Its three sets of scalar fields that you are going to have to combine into a vector field in VisIt using expressions. Go make this.

Your first task is to open up your dataset using just the cli. Maybe there is a function that will do this for you. Maybe its called `OpenDatabase(databaseName)`. Maybe `databaseName` will be set to something like `"/home/projects/TutorialData/myfilename_*.vtk database"` Make sure that your database has opened without errors. Now would be a good time to make a script. In your favorite text editor (*cough *cough vim *cough), open up a new document and call it whatever you want. Make sure to add `.py` to the end. At the top of your new script put `"Source(/path/to/script.py)"` This will help because you can't copy and paste into the cli normally because its stupid. Instead, you have to highlight what you want to copy and paste it by clicking with the middle mouse button (This magical thing you just did is a very simple copy and paste. highlighting text will copy it into a buffer, and clicking the middle mouse pastes it to where the cursor is. Learn to use this tactic, as it's going to be very handy) This helps when you need to rerun things and don't want to type a full path all the time.

Now that that's done, delete all the plots. You'll need this because otherwise if you rerun your script, you'll end up with multiple plots. Next we want to make a plot. For now we want a streamline plot. Also make sure to add all the settings. Next you can do things like setup the settings for window saving, change the background color to your favorite color, remove the bounding box, set your initial view angle, etc. There are functions that will do this. Usually one of them includes the words `Get` and `Attributes`. You can probably

find it in the Python interface manual. Now, you can begin coding the movie making loop. You'll want a for loop that iterates through each frame once (hint: `range(GetNSteps())`). For each iteration, you'll set the timestep, make the changes that you want (i.e. add new streamline seeds, change the view angle, etc.), and call `SaveWindow()`. For now, just update the timestep and call `SaveWindow()`. With the above done, you should have a basic movie script. You can run your script in the cli using `"Source(/path/to/script.py)"`. To make things more interesting, you are now going to add code that will update the streamline plot so that the streamline seeds will change each timestep. This is pretty important for actual movies since we often use particles that move through time as seeds. The seed points are given to us (if they are particles) or we can generate them ourselves (we have to process them into .txt files though). This time, just generate some seeds. Edit your code that you used to visualize a rotating sphere to just give a txt file with some rotating points. Or, you could make moving points in some other way. It doesn't matter too much for the purpose of this exercise. To see a good amount of streamlines, I'd make it so there are about maybe 10 points total on the sphere (or whatever), but that's just a guesstimate. You can do more or less depending on how you feel it will look. You can also always go back and change it. With this data and updated code, make a new movie where the streamline seeds are updated at each timestep. Now add code to make the camera pan across your data set. probably use something like `cubicevalspline` to help you with this. 9 With that done, you've done most of the fancy things you need to do in an ordinary movie.

8 Blue Waters

The size of our simulation data is on the order of terabytes. Thus it requires a mighty machine to both store that data, and parse through it quickly. This is where Blue Waters comes in. Each simulation is a couple hundred frames long, and each frame takes about 20 minutes to create. Blue Waters allows us generate these frames in a matter of minutes, rather than days. This is done using parallel programming. Read the getting started guide on the blue waters page <https://bluewaters.ncsa.illinois.edu/getting-started>.

You are now going to edit your movie script so that it can be run on Blue Waters. Hopefully by now, you guys have accounts on Blue Waters. If not you can just borrow anyone else's account for the purpose of this tutorial.

The first step is moving your data to Blue Waters. There are two ways to do this. One is to use a terminal command such as `scp` or `rsync`. The other is a program called `globus online`.

`Globus online` should be used for transferring large amounts of data because that's how the BW guys want you to do it. You can use the terminal for smaller files such as scripts, seeds, and .txt files. Even though your data set is not very large, I'm going to have you transfer the data using `globus online`. Ask me how to do that when you get to this point. Then, when you are done with that `scp` or `rsync` the rest of your files over. Nodes on blue waters can be thought of as

individual computers. There are two types of nodes found on any computing system of this type. The first is a login node. These are the nodes which you automatically ssh onto. The second are the compute nodes.

You can't log on to these nodes but you can use a command called `qsub` to submit jobs. The jobs are submitted to a queue and run when its your turn. The default compute node has something like 32GB of RAM split up between 16 processor cores. These nodes are all connected with a very fast ethernet connection so they can send data to each other really fast. They are also all connected to the same data on disk as well. change your `.py` script so that the files you access correspond to where you data is on bw. Then go here <https://bluewaters.ncsa.illinois.edu/visit> and read the section on visit's command line interface and batch mode. You can use the sample `.pbs` file they give you. With that file in hand, call `"qsub myfile.pbs"` after editing the file.

You can check on your submission with `"qstat -u username"` (or ask one of the older guys for their `~/.bashrc` file). If that works, now try and "parallelize" your code. More specifically, you want to make it so that you can split up the frame making work amongst multiple computers. For example, say your dataset has 200 timesteps. You have 4 computers. That means you'll probably want to give each computer 50 timesteps to make pictures from. Make adjustments in your code so that when given the total number of computers and the computer number (we'll call it rank), you'll successfully split up the work and output frames for you movie. In our example, we have 4 computers total. One computer will have rank 0, another will have rank 1, etc. However, don't hardcode these numbers directly. Keep them as variables which you can pass to your script. To access them in your code, use `sys.argv`. If these concepts don't make sense, feel free to ask me for clarification.

8.1 Selected Exercises:

1. Log into Blue Waters using `"ssh -X 'username'@bw.ncsa.illinois.edu"`
2. Open Visit using `"module load visit; visit"`. Check the settings to ensure that the Delaunay operator is loaded. Close VisIt.
3. Make a batch script following the recipe found online, and submit it to the blue waters job queue.
4. Bask in the glory of your life because you're using a heavenly supercomputer.

9 Particle Picker

We need particles to create a seeding points for magnetic fieldlines and to create particle movies. You will get a raw data file for particle position at all time, usually named as `"bhns-particles.mon"` or `"particles.mon"`. Our task is to choose

some useful particles, and convert the file into a format that is readable in VisIt. This can be done by using the particle picker module.

You will need to transfer the abid bot to your Blue Waters. This is the ultimate code that you will be working with most of the time. For now we will skip most of the features and focus on particle picker module. The module is located in `<path to your abid bot>/bin/particle_code/`. But there are a few steps you need to do before getting there.

1. Ask for the "particles.mon" file from previous REU members. Then copy it to `<path to your abid bot>/h5data/`. Note that you need to rename it to "particles.mon" if it is named differently.
2. Change parameters inside the file `<path to your abid bot>/params`. We are going to work with particle seeds, so set `fields=true`, `particleSeeds=true`, and `updateParticleMon=true`. Other parameters will be different from case to case. Ask others if you are stuck here for too long.
3. Now setup the parameters. First, enter `". params"` in command line. This will save parameters that you specified in params on your current shell. Next, normally you will run the `setup.sh` script to set everything. But since we are only doing particle picker, you can run abid bot partially by going into bin folder and run `". setup_seeds.sh"`. This step will take considerably long time, usually about 1+ min per GB of your particles.mon file. After this is done, go to `<path to your abid bot>/bin/particle_code/` and check "dat" folder. You should have a bunch of files with the name like "00012.34567000000.dat" and the name of your first file should be "(firstTime).dat". If you get this then Cheers!, you've finished the most time consuming part of particle picker. Set `updateParticleMon=false` in params to skip this part in future runs, as it is only necessary to do so once.

Now go to `<path to your abid bot>/bin/particle_code/`. There are three files involving particle picker: `particlePicker.py`, `particlePicker-tracer.py`, and `particlePickerModule.py`. Do not edit the `particlePickerModule.py` unless you know what you are doing. By all means, I encourage you to take a look inside and try to understand the module. The file you will be editing is either `particlePicker.py` or `particlePicker-tracer.py`. They are exactly the same except that the tracer one save files in .3d format instead of .txt. These script will pick particles based on your criteria. There are currently two methods to choose particles: `findInVolume` and `nearestNeighbor`. Details for these methods are in `particlePickerModule.py`.

Your task for this part is to make a simple particle tracer movie (only particle, without density and fieldlines). If you don't know what it is, please see http://research.physics.illinois.edu/cta/movies/NSNS_high_res/particles_ext.html. You will have to edit `particlePicker-tracer.py`, set `particleTracer=true` in params, run `". params"`, and run `". setup_seeds.py"`. If it is working correctly, you will get `trace*.3d` files inside `tracel` folder. Copy these files to your local machine

and open them in VisIt (You may need to rename them so that VisIt can see them as a database). Then plot them using Pseudocolor plot. I will let you play with VisIt and figure how to create a movie from here.

10 The Final Task

It's now time for you to get your hands on some real data. Ask one of the senior members to provide you with an old case to work with. This step can only be done on Blue Waters, as the data is too big for our local machines. While the data is transferring on your BW, transfer the file `abid_bot.tar.gz` file there if you have not done so. Go through the package (hint: `README` is a good place to start) and try to understand what's going on in different parts. Use this code to make a movie. This movie will be your certificate that you've successfully completed the tutorial.

11 Bonus: Another Dimension

Let's redo the Final Task, but this time in 3D! Before beginning, ask a senior member to play a sample 3D movie created with VisIt.

11.1 3D Movie exercise:

1. The *.png generating process is almost identical to the previous section, and you may make the 3D movie using the same data. In `run_movie_ranks.py` (found in `abid_bot/bin/bw_many_folder_scripts/`), simply enable the "stereo" `SaveWindowAttribute`.
2. You may want to change the view to one that works better with 3D. The file containing the view attributes is specified in `params`, and you can change this to your own view .xml file. One attribute you can try changing is the `eyeAngle`. Hint: Open a single frame of your movie in VisIt, and change the view to your liking. Then you can export the `View3DAttributes` to an .xml using the `SaveAttribute()` function. See the VisIt Python Interface Manual for more details.
3. Use the `abid_bot` code to generate *.png images as before.
4. To turn your images into a .mp4 movie, you will need to run two scripts: `merge_3D.sh`, and then `Movie_maker_3D.sh`. You can transfer your `movies/` folder from Blue Waters to your computer if you'd like to run the scripts locally. Before running the scripts look them over and make necessary changes. If you don't have the scripts ask a group member.
5. You should now have a .mp4 movie. To play it, you can use the command "mplayer filename". To actually see the movie in 3D, you'll need special equipment- ask a senior member for help. You now have a 3D movie worthy of IMAX.

12 Final Thoughts

Congratulations! While there are a couple of things that you will need to do, such as working with kdenlive and editing the website, those things are best taught by joining in the next time the group as a whole needs to use those tools. You should now be ready to jump right in to the group and start making science!